

HPD Analysis (No MAP)-Copy1

September 1, 2024

1 Houston Police Department Elementary Data Analysis

1.0.1 Noah Choate

1.1 Data Information

1.1.1 About the Data

For the purpose of this assignment, the dataset that will be analyzed is the National Incident-Based Reporting System's (NIBRS) dataset in the Houston area. The NIBRS is responsible for capturing characteristics of incident based data, such as the date and time, type of incident, and location. The dataset provided details all criminal incidents in the Houston area in 2023.

The dataset consists of 16 columns and 249737 rows. In the analysis, I will begin by exploring the trends of when incidents occur by months / seasons of the year, and hours of the day in attempts to see if there are any trends. Following that by detailing the frequency of when the incidents happen in the top zip-codes, along with the top descriptions (Theft from motor vehicle, etc.). Concluding the analysis will be a spatial analysis utilizing k-clustering that will aim to find common areas of where incidents occur, beside a concluding statement to finish the report.

1.1.2 Loading the Libraries / Data

To begin, I first loaded libraries that will aid me in conducting my data analysis. Pandas and Numpy in order to help manipulate the dataframe to conduct a fair analysis as well as matplotlib and seaborn to produce visuals. I also had imported a Linear Regression function in order to help my analysis on the frequency of crimes throughout the year. In order to do spatial analysis, I had imported KMeans for clustering and folium to visualize the clusters on a map.

```
[10]: # Import Necessary Libraries
import pandas as pd
import geopandas as gpd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
sns.set(style = "darkgrid")
from sklearn.linear_model import LinearRegression
from sklearn.cluster import KMeans
import folium
```

```
# Read dataframe in
df = pd.read_excel(r"C:\Users\nahc\Documents\Algorithms and Data Structures\data\NIBRSPublicView2023.xlsx")

# Preview of df
print(df.head())
```

	Incident	RMSOccurrenceDate	RMSOccurrenceHour	NIBRSClass	\
0	2423	2023-01-01	0	90D	
1	2823	2023-01-01	0	520	
2	2823	2023-01-01	0	35A	
3	4623	2023-01-01	0	220	
4	6723	2023-01-01	0	13B	

	NIBRSDescription	OffenseCount	Beat	\
0	Driving under the influence	1	5F20	
1	Weapon law violations	1	10H50	
2	Drug, narcotic violations	1	10H50	
3	Burglary, Breaking and Entering	1	5F30	
4	Simple assault	1	14D20	

	Premise	StreetNo	StreetName	StreetType	Suffix	\
0	Highway, Road, Street, Alley	1760	CAMPBELL	RD	NaN	
1	Residence, Home (Includes Apartment)	3100	ANITA	ST	NaN	
2	Residence, Home (Includes Apartment)	3100	ANITA	ST	NaN	
3	Bank, Savings & Loan	5253	HOLLISTER	ST	NaN	
4	Residence, Home (Includes Apartment)	4119	BARBERRY	DR	NaN	

	City	ZIPCode	MapLongitude	MapLatitude
0	HOUSTON	77080	-95.516391	29.802506
1	HOUSTON	77004	-95.357674	29.732490
2	HOUSTON	77004	-95.357674	29.732490
3	HOUSTON	77040	-95.506306	29.842607
4	HOUSTON	77051	-95.367319	29.651655

1.1.3 Preparing / Cleaning the Data

Due to the nature of the raw dataframe given by the NIBRS, it will be necessary to perform some manipulation on the dataframe in order to improve the accuracy of my data analysis. One of the top priorities when cleaning data is to remove all of the missing data in the dataframe. Because our data is mainly categorical and independent events, it does not make sense to speculate/ predict data. One glaring issue is The “Suffix” column in the dataframe. This column was used to clarify what road an incident occurred, (Ex: Shepherd Drive “N”) however since many roads do not utilize a suffix, removing this column is necessary. If we do not remove this column prior to removing all missing data, over 75% of our data will be removed, and therefore the report will not be as accurate as it could be. We will then locate where missing values are, and drop them. Additionally it is vital to ensure all of our datatypes in our dataframe are the appropriate type in order to perform

analysis with them.

```
[2]: # Remove Suffix col, will remove too much data if kept upon cleaning. Redundant
df.drop(columns = ['Suffix'], inplace = True)

# Confirm the correct data types
print(df.dtypes)

# locate this missing values
missing_values = df.isnull().sum()
print(missing_values)

# Because the dataset is categorical, remove all rows with missing values
df_dropNA = df.dropna()
print(df)
```

```
Incident          int64
RMSOccurrenceDate  datetime64[ns]
RMSOccurrenceHour  int64
NIBRSClass         object
NIBRSDescription   object
OffenseCount       int64
Beat              object
Premise            object
StreetNo           object
StreetName         object
StreetType         object
City              object
ZIPCode           object
MapLongitude       float64
MapLatitude        float64
dtype: object
Incident          0
RMSOccurrenceDate  0
RMSOccurrenceHour  0
NIBRSClass         0
NIBRSDescription   0
OffenseCount       0
Beat              193
Premise            0
StreetNo           860
StreetName         0
StreetType         18276
City              1
ZIPCode           3393
MapLongitude       3190
MapLatitude        3190
dtype: int64
```

	Incident	RMSOccurrenceDate	RMSOccurrenceHour	NIBRSClass	\
0	2423	2023-01-01	0	90D	
1	2823	2023-01-01	0	520	
2	2823	2023-01-01	0	35A	
3	4623	2023-01-01	0	220	
4	6723	2023-01-01	0	13B	
...	
249733	1601024	2023-12-31	23	240	
249734	185927523	2023-12-31	23	13A	
249735	185940723	2023-12-31	23	35A	
249736	185940723	2023-12-31	23	90D	
249737	185962423	2023-12-31	23	290	

	NIBRSDescription	OffenseCount	Beat	\
0	Driving under the influence	1	5F20	
1	Weapon law violations	1	10H50	
2	Drug, narcotic violations	1	10H50	
3	Burglary, Breaking and Entering	1	5F30	
4	Simple assault	1	14D20	
...	
249733	Motor vehicle theft	0	5F20	
249734	Aggravated Assault	1	11H50	
249735	Drug, narcotic violations	1	1A10	
249736	Driving under the influence	1	1A10	
249737	Destruction, damage, vandalism	1	16E10	

	Premise	StreetNo	StreetName	\
0	Highway, Road, Street, Alley	1760	CAMPBELL	
1	Residence, Home (Includes Apartment)	3100	ANITA	
2	Residence, Home (Includes Apartment)	3100	ANITA	
3	Bank, Savings & Loan	5253	HOLLISTER	
4	Residence, Home (Includes Apartment)	4119	BARBERRY	
...	
249733	Other, Unknown	7930	BLANKENSHIP	
249734	Residence, Home (Includes Apartment)	4800	ALLENDAL	
249735	Highway, Road, Street, Alley	1800	PEASE	
249736	Highway, Road, Street, Alley	1800	PEASE	
249737	Residence, Home (Includes Apartment)	13351	CITY PARK CENTRAL	

	StreetType	City	ZIPCode	MapLongitude	MapLatitude
0	RD	HOUSTON	77080	-95.516391	29.802506
1	ST	HOUSTON	77004	-95.357674	29.732490
2	ST	HOUSTON	77004	-95.357674	29.732490
3	ST	HOUSTON	77040	-95.506306	29.842607
4	DR	HOUSTON	77051	-95.367319	29.651655
...
249733	DR	HOUSTON	77055	-95.488949	29.819928
249734	RD	HOUSTON	77017	-95.250621	29.681868

249735	ST	HOUSTON	77003	-95.362190	29.747570
249736	ST	HOUSTON	77003	-95.362190	29.747570
249737	LN	HOUSTON	77047	-95.384408	29.614650

[249738 rows x 15 columns]

To ensure the analysis being performed is an accurate representation of the total dataframe, it is vital that we minimize the percentage of values dropped. A percentage of values dropped being below 10% in a dataframe containing roughly 250,000 entries will still provide an accurate analysis.

```
[3]: # Clarify how much data was removed
before_length = len(df)
print(f'Length before dropping NA values: {len(df)}')

after_length = len(df_dropNA)
print(f'Length after dropping NA values: {len(df_dropNA)}')

percent_decrease = (abs(before_length - after_length) / before_length) * 100
print(f'Percentage of values dropped: {percent_decrease}%')
```

Length before dropping NA values: 249738
Length after dropping NA values: 226988
Percentage of values dropped: 9.109546805051693%

Furthermore, it would also be noteworthy to understand how many unique instances there are for the columns in the dataframe I will be doing an analysis on. I chose ZIP-code over City as there are more zip-codes listed, and therefore will be a more accurate display of percentages given by the analysis.

```
[4]: # Zipcode
unique_zips = df['ZIPCode'].nunique()
print(f'Number of Unique Zip-Codes Listed: {unique_zips}')
```

```
# Description
unique_descriptions = df['NIBRSDescription'].nunique()
print(f'Number of Unique Descriptions Listed: {unique_descriptions}')
```

```
# NIBRSClass
unique_class = df['NIBRSClass'].nunique()
print(f'Number of Unique NIBRS Classes Listed: {unique_class}')
```

```
# Date
unique_date = df['RMSOccurrenceDate'].nunique()
print(f'Number of Unique Dates Listed: {unique_date}')
```

```
# Lat/Long
unique_lat = df['MapLatitude'].nunique()
print(f'Number of Unique Latitudes / Longitudes Listed: {unique_lat}')
```

Number of Unique Zip-Codes Listed: 212

Number of Unique Cities Listed: 141
Number of Unique Descriptions Listed: 61
Number of Unique Beats Listed: 129
Number of Unique NIBRS Classes Listed: 61
Number of Unique Dates Listed: 365
Number of Unique Latitudes / Longitudes Listed: 69580

1.2 Descriptive Statistics and Visualization

To begin the analysis we will look a time-series analysis regarding seasonal and trends of crime, followed by a look into the distribution of incidents by the hour of the day. Following the time analysis, will be an examination of the frequency of the top NIBRS Descriptions of incidents, as well as the ZIPCodes. Finalizing the report will be a spatial analysis to identify hotspots in the Houston area of crime using k-clustering.

1.2.1 Linear Regression Time-Series Analysis

```
[5]: # Prepare Data for Linear Regression on Dates / Seasons

# Season Work
df['Month'] = df['RMSOccurrenceDate'].dt.month

seasons = {
    'Winter': [12, 1, 2],
    'Spring': [3, 4, 5],
    'Summer': [6, 7, 8],
    'Fall': [9, 10, 11]
}

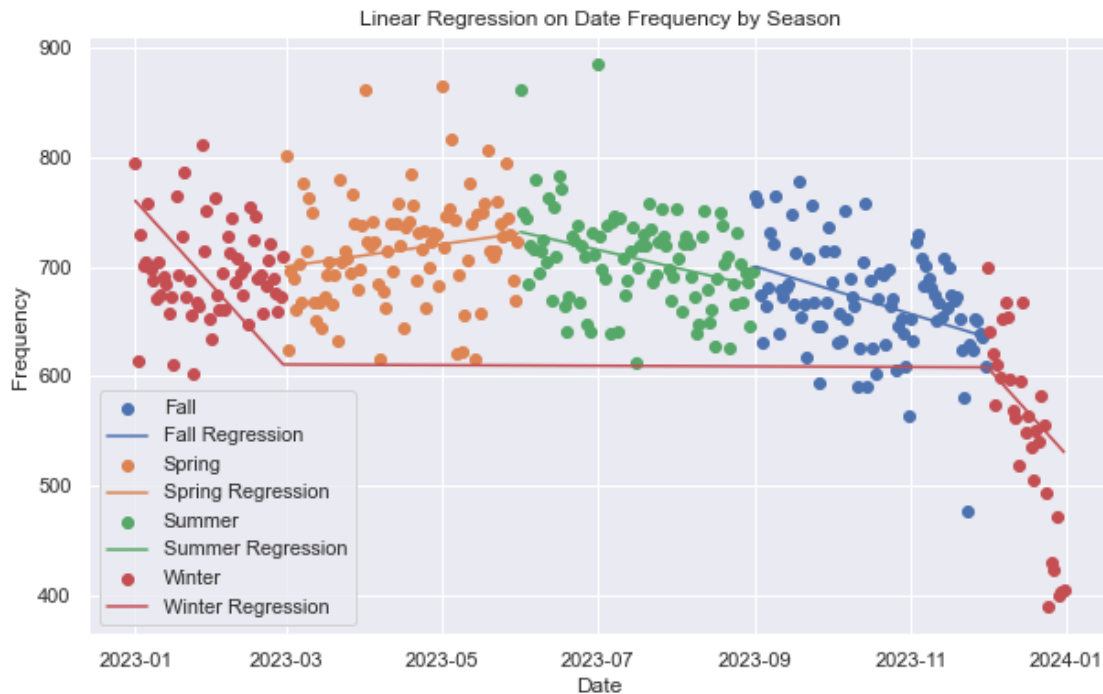
# Map months to seasons
df['Season'] = df['Month'].apply(lambda x: next(season for season, months in
↪seasons.items() if x in months))

season_frequency = df.groupby(['Season', 'RMSOccurrenceDate']).size().
↪reset_index(name='frequency')

# Plot linear regression for each season
plt.figure(figsize=(10, 6))
for season, data in season_frequency.groupby('Season'):
    X = data.index.values.reshape(-1, 1)
    y = data['frequency']
    model = LinearRegression()
    model.fit(X, y)
    y_pred = model.predict(X)
    plt.scatter(data['RMSOccurrenceDate'], data['frequency'], label=season)
    plt.plot(data['RMSOccurrenceDate'], y_pred, label=f'{season} Regression')

plt.xlabel('Date')
```

```
plt.ylabel('Frequency')
plt.title('Linear Regression on Date Frequency by Season')
plt.legend()
plt.show()
```



The plot above is displaying the seasonal trends of the frequency of dates incidents occur in. According to the data, it is common for the Houston area to begin the year slowly creeping from around the 700 incidents per day to roughly 725 incidents during the peak around May. Slowly towards the end of the year we can expect incidents to drop in October from roughly 700 towards 650, only to drop 625 in early December. Oddly, in late December there is a sharp decrease that finishes the year around 550 incidents. If we speculate this graph to repeat, then we can expect a sharp rise during the new year. Utilizing previous years' data would help this prediction, however for the purpose of this elementary data analysis that will not be necessary. This information is important as it could be a reasoning to re-tool / re-manage our police departments. For example, during December when incidents reach a low, maybe shuffle around the police department to have less people patrolling during December. Conversely, slightly increase the number of police on patrol around May, when incidents are at a yearly-high.

1.2.2 Distribution Time-Series Analysis

```
[6]: # Prepare data for visualization of distribution of incidents per hour of the day
df['RMSOccurrenceHour'] = pd.to_datetime(df['RMSOccurrenceHour'], format='%H')

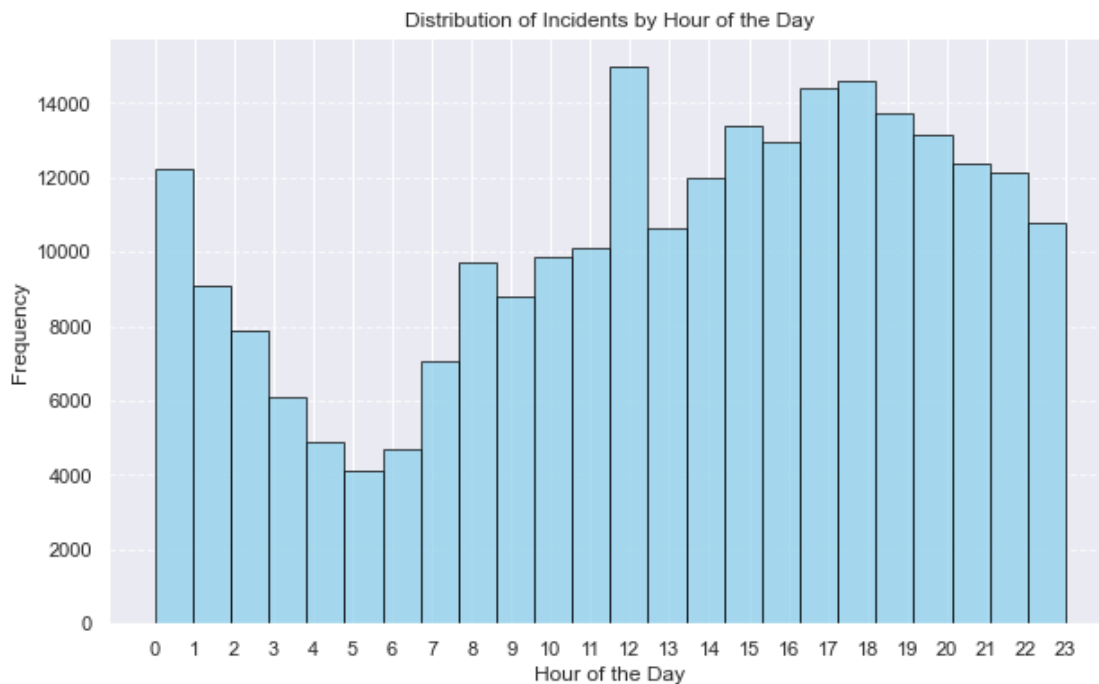
# Extract hour
```

```

df['Hour'] = df['RMSOccurrenceHour'].dt.hour

# Histogram to show the distribution of incidents by hour
plt.figure(figsize=(10, 6))
plt.hist(df['Hour'], bins=24, color='skyblue', edgecolor='black', alpha=0.7)
plt.title('Distribution of Incidents by Hour of the Day')
plt.xlabel('Hour of the Day')
plt.ylabel('Frequency')
plt.xticks(range(24)) # Set ticks for each hour
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.show()

```



The histogram above displays the distribution of incidents and the hour of the day they occur. From the data, we can observe that most incidents occur at noon. Throughout the day, you can typically expect most incidents are likely to occur around 4-6 PM. This information can be useful as it can show when our police department can be utilized most efficiently on the street. Because of the analysis above, it is proven that obviously it would be more beneficial to have the majority of our officers patrolling from noon to 6 PM rather than midnight to 6 AM.

1.2.3 Data Frequency Analysis

Pie Chart

```

[7]: # Top 10 most common ZIP-Code occurrences
common_zip = df['ZIPCode'].value_counts()
ten_common_zips = common_zip.head(10)

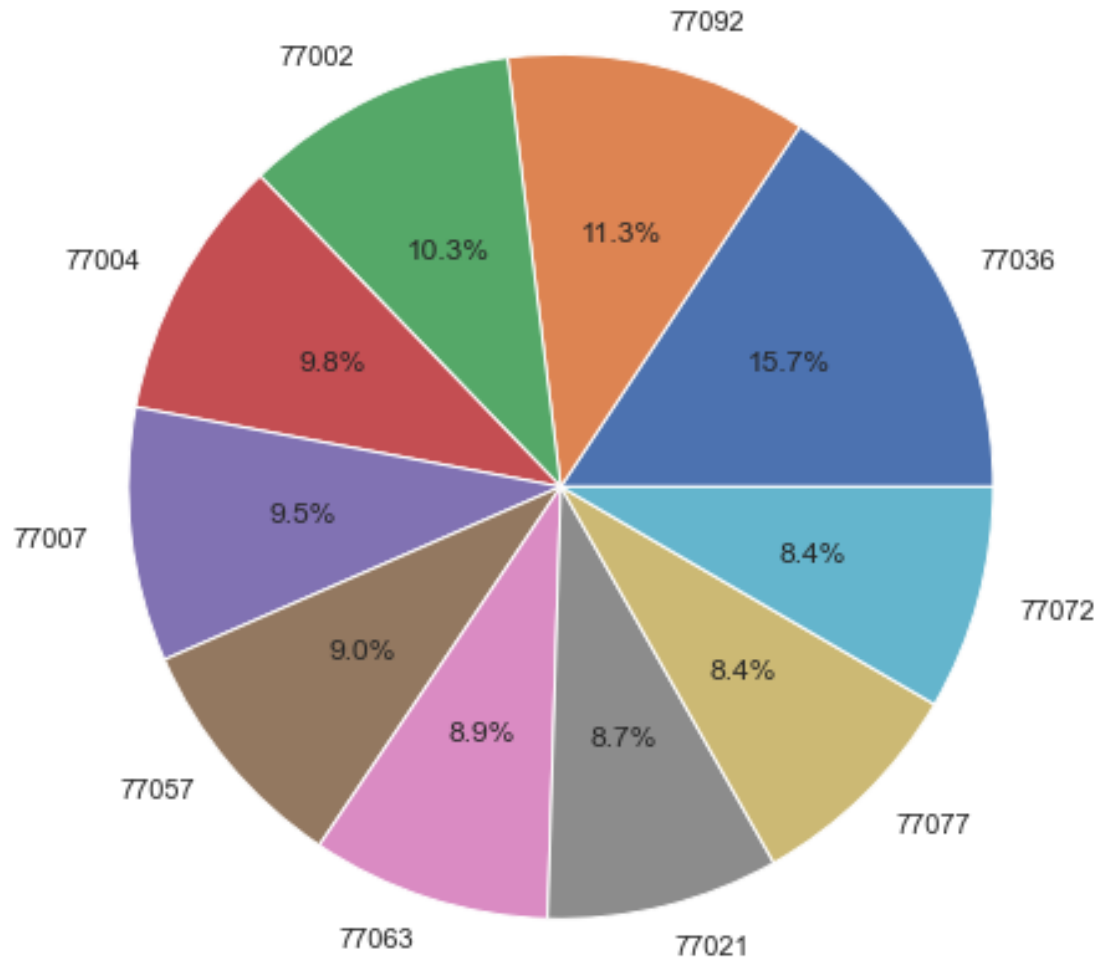
```

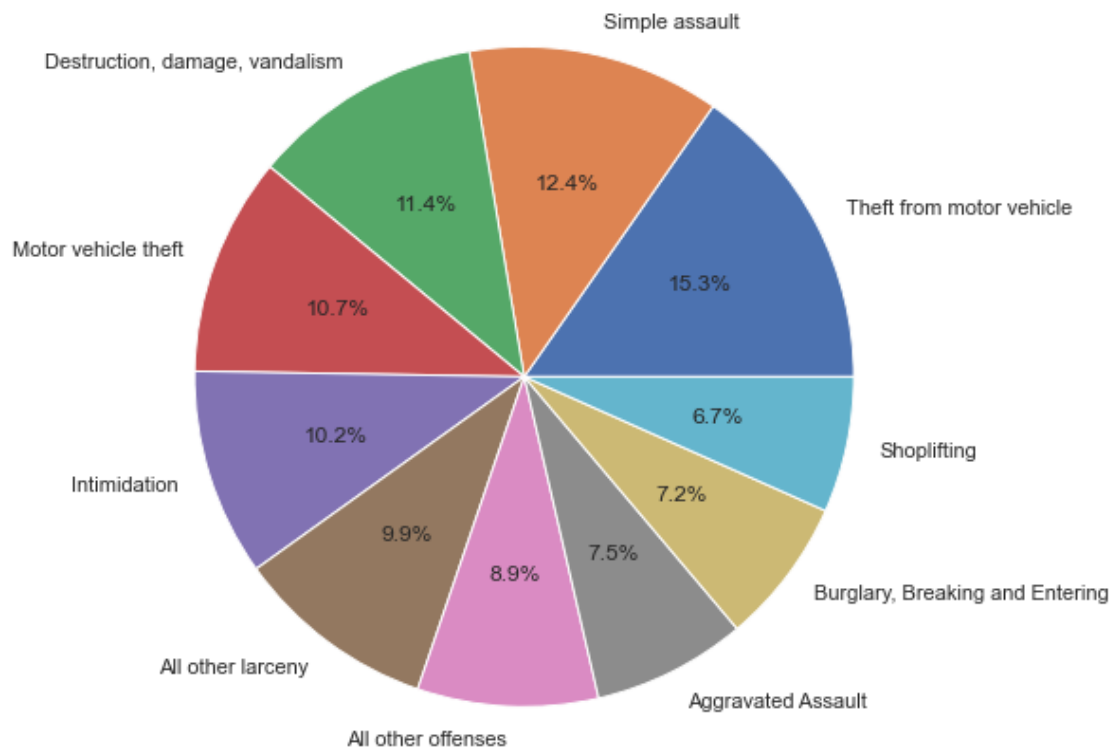


```
plt.figure(figsize=(8,8))
plt.pie(ten_common_zips, labels = ten_common_zips.index, autopct = '%1.1f%%')

# Top 10 most common Offenses
common_offenses = df['NIBRSDescription'].value_counts()
ten_common_offenses = common_offenses.head(10)
plt.figure(figsize=(8,8))
plt.pie(ten_common_offenses, labels = ten_common_offenses.index, autopct = '%1.1f%%')
```

```
[7]: ([<matplotlib.patches.Wedge at 0x27f4fe2c850>,
      <matplotlib.patches.Wedge at 0x27f4fe2cfa0>,
      <matplotlib.patches.Wedge at 0x27f50214700>,
      <matplotlib.patches.Wedge at 0x27f50214e20>,
      <matplotlib.patches.Wedge at 0x27f50218580>,
      <matplotlib.patches.Wedge at 0x27f50218ca0>,
      <matplotlib.patches.Wedge at 0x27f50212400>,
      <matplotlib.patches.Wedge at 0x27f50212b20>,
      <matplotlib.patches.Wedge at 0x27f50205280>,
      <matplotlib.patches.Wedge at 0x27f502059a0>],
      [Text(0.9755569626079338, 0.5082210274154176, 'Theft from motor vehicle'),
       Text(0.2414622532201113, 1.0731709930248148, 'Simple assault'),
       Text(-0.5507505908462679, 0.9521941958878384, 'Destruction, damage,
vandalism'),
       Text(-1.0320964793133083, 0.38049554187805396, 'Motor vehicle theft'),
       Text(-1.0500947576358097, -0.3275683134642152, 'Intimidation'),
       Text(-0.6562450771013095, -0.8828037147522072, 'All other larceny'),
       Text(-0.05572681264412226, -1.0985875123778384, 'All other offenses'),
       Text(0.49323525991956685, -0.9832186828839643, 'Aggravated Assault'),
       Text(0.8799881555257706, -0.6600157923370867, 'Burglary, Breaking and
Entering'),
       Text(1.076022474741318, -0.2284198630845616, 'Shoplifting')],
      [Text(0.5321219796043275, 0.2772114694993187, '15.3%'),
       Text(0.13170668357460616, 0.5853659961953535, '12.4%'),
       Text(-0.3004094131888733, 0.5193786523024573, '11.4%'),
       Text(-0.5629617159890772, 0.20754302284257486, '10.7%'),
       Text(-0.5727789587104416, -0.17867362552593558, '10.2%'),
       Text(-0.3579518602370779, -0.4815292989557493, '9.9%'),
       Text(-0.03039644326043032, -0.5992295522060935, '8.9%'),
       Text(0.26903741450158186, -0.5363010997548896, '7.5%'),
       Text(0.479993539377693, -0.3600086140020473, '7.2%'),
       Text(0.5869213498589007, -0.12459265259157905, '6.7%')])
```





Table

```
[8]: # Calculate frequency and percentage of ZIP codes
zip_count = df['ZIPCode'].value_counts()
total_times = zip_count.sum()
zip_percentage = (zip_count / total_times) * 100

# Create df for ZIP codes
zip_table = pd.DataFrame({'ZIP-Code': zip_count.index,
                          'Frequency': zip_count.values,
                          'Percentage': zip_percentage.values})

print("ZIP Code Table:")
print(zip_table)

# Calculate frequency and percentage of NIBRS descriptions
desc_count = df['NIBRSDescription'].value_counts()
total_descs = desc_count.sum()
desc_percentage = (desc_count / total_descs) * 100

# Create df for NIBRS descriptions
```

```
desc_table = pd.DataFrame({'Description Type': desc_count.index,
                           'Frequency': desc_count.values,
                           'Percentage': desc_percentage.values})

print("\nNIBRS Description Table:")
print(desc_table)
```

ZIP Code Table:

	ZIP-Code	Frequency	Percentage
0	77036	9127	3.704967
1	77092	6576	2.669427
2	77002	5976	2.425866
3	77004	5718	2.321135
4	77007	5532	2.245631
..
207	75050	1	0.000406
208	77045-0000	1	0.000406
209	91406	1	0.000406
210	77001	1	0.000406
211	77060-1668	1	0.000406

[212 rows x 3 columns]

NIBRS Description Table:

	Description Type	Frequency	Percentage
0	Theft from motor vehicle	28234	11.305448
1	Simple assault	22865	9.155595
2	Destruction, damage, vandalism	20973	8.398001
3	Motor vehicle theft	19779	7.919900
4	Intimidation	18761	7.512273
..
56	Incest	5	0.002002
57	Human Trafficking/Involuntary Servitude	2	0.000801
58	Negligent manslaughter	2	0.000801
59	Runaway	2	0.000801
60	Peeping tom	1	0.000400

[61 rows x 3 columns]

With categorical data, it is important to note the frequency and relative distribution of each element. In this scenario, it is important to analyze where these incidents are likely to occur and what incidents occur. To begin, a pie-chart visualizing the the zip-codes where most incidents occur, followed by a in-depth table. One thing to note is that the pie-chart display percentages in proportion to the top-10 zip-codes where incidents occur, while the table displays percentages in regards to all zip-codes. Because all of these zip-codes are in inner-Houston and Harris County, it was important to narrow the location search down to zip-code (for now). This is important as again, knowing what zipcodes most incidents occur in gives justification to toughen up patrol out in those areas in hopes to prevent more incidents from occurring. For example, we can see that not

only the most- but double the incidents occur in 77036 than in 77077. With this knowledge, there is concrete evidence that more patrolling needs to be done in those areas than say 77001 where only 1 incident had occurred last year.

Our second pie-chart and table combination details what NIBRS description type each incident had received. Again, the pie-chart shows the percentages relative to the top-10, while the tables give the percentages regarding all description types. This visualization is important as it shows what category Houston needs help with the most in regards to illegal incidents. From the table, we can tell that roughly 11% of all incidents are Theft from a motor vehicle. With this knowledge, we can train new police officers how to deal with the most common situation they will likely face on the job, as well as measure to prevent these from occurring. If we know theft from a motor vehicle is common, then we can encourage local businesses to install cameras in their parking lots to prevent customer theft, tell citizens to keep valuables out of sight. If we know simple assault adds up to roughly 9% of all incidents in the Houston area, we can inform citizens of the risk and advise them to carry modes of self-defense whether that be pepper-spray or a handgun.

1.2.4 Data Spatial Analysis with K-Clustering

```
[1]: # Code included in .html / .ipynb file. Not available on PDFs
```

Spatial analysis utilizing K-Clustering is helpful to find hotspots of incidents occurring. In this scenario, a K-Value of 13 is chosen for the 13 divisions in Houston. While each division is not labeled on the map, the 13 clusters are intended to show a possible re-mapping of the divisions in terms of efficiency regarding incidents rather geographic location. This is important as in potential districts converging, diverging, dissolves, or remapping, this can be an outline for a beneficial way to remap our districts to protect the community. This could also help districts relocate and remanage current positions of where officers patrol or how many officers patrol a designated area. Alternatively, this analysis helps position our police force to be prepared for the harm that comes in our communities' way.

1.3 Conclusion

Data Analysis can be highly beneficial when it comes to protecting our communities. With analysis' such as regression to check on patterns in yearly/ daily trends, K-Clustering to see what areas are highly dense and even having the frequency of incidents occurring can provide helpful stepping stones in organizing our police force. Through the use of the data analysis, we have discovered that most incidents peak during the summertime and dramatically drop during December to quickly pick up a roughly 700 incidents a month until the cycle repeats. We have identified the most dangerous zip-codes and descriptions of where incidents occurs to help locate and inform our officers on where to go and what challenges they will likely face performing their job. Lastly, we have performed K-Clustering utilizing $K = 13$ (The number of divisions in Houston) to help strategically place our officers in position for the best possible performance. While this analysis only consists of data in 2023, the more and more data we insert into this analysis the more beneficial it would provide. It would prove worthwhile to revisit this dataset and add new and old entries to help us in the future, as the more data we have the more accurate our numbers become.